

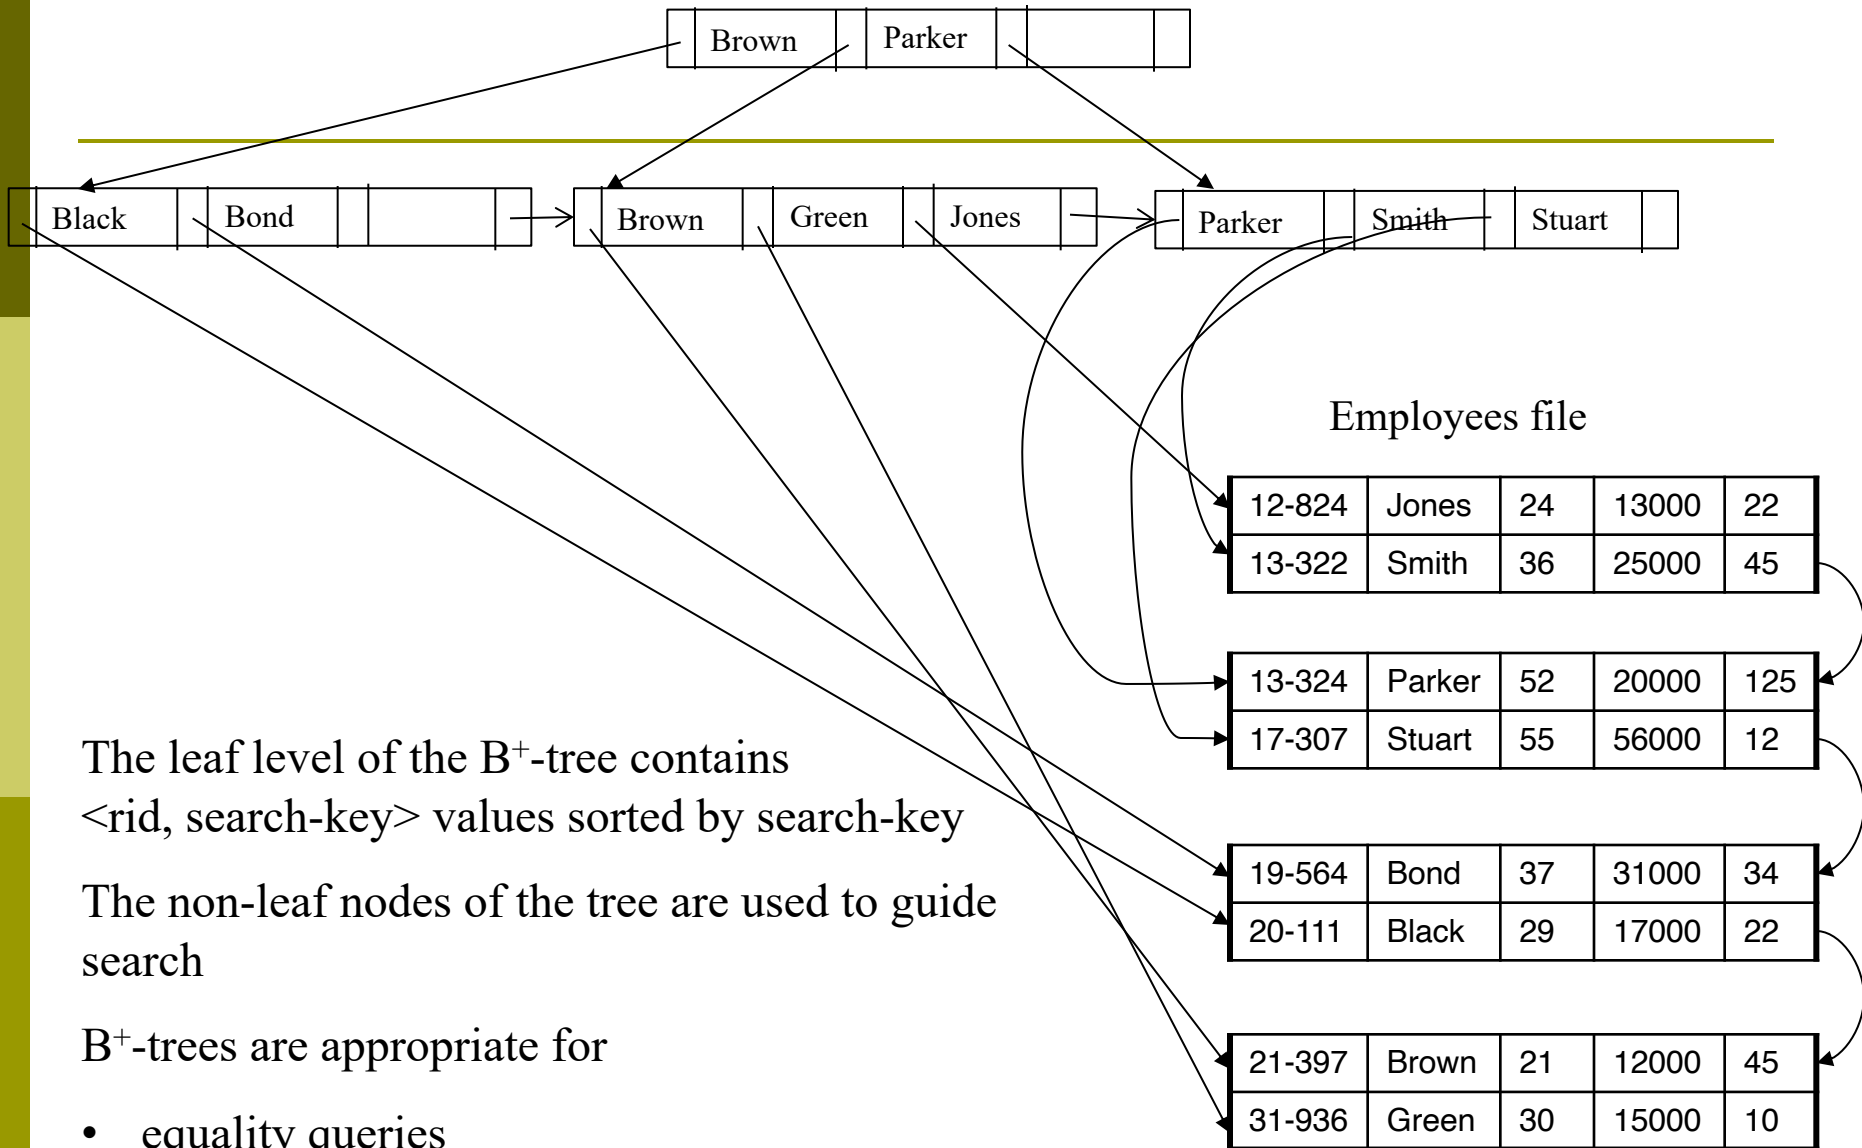
Adaptive Indexing

- ❑ Why adaptive indexing?
- ❑ Standard database cracking
- ❑ Adaptive merging
- ❑ Stochastic cracking
- ❑ Coarse granular index
- ❑ Multidimensional adaptive indexing

Background: Database Indexing

- **Indexing** is a fundamental approach in database systems that facilitates search
- **Objective**: given a search-key value q for attribute A of a DB table T , find fast the record(s) r in T for which $r.A = q$
 - **Range query**: find records s.t. $q_1 \leq r.A \leq q_2$
- An index is a table of **<search-key, rid>** pairs, organized in a data structure
 - search-key is a value of A in table T
 - rid is an identifier or pointer to a record (or set of records) in T for which $A = \text{search-key}$

B+-Tree Index: Example



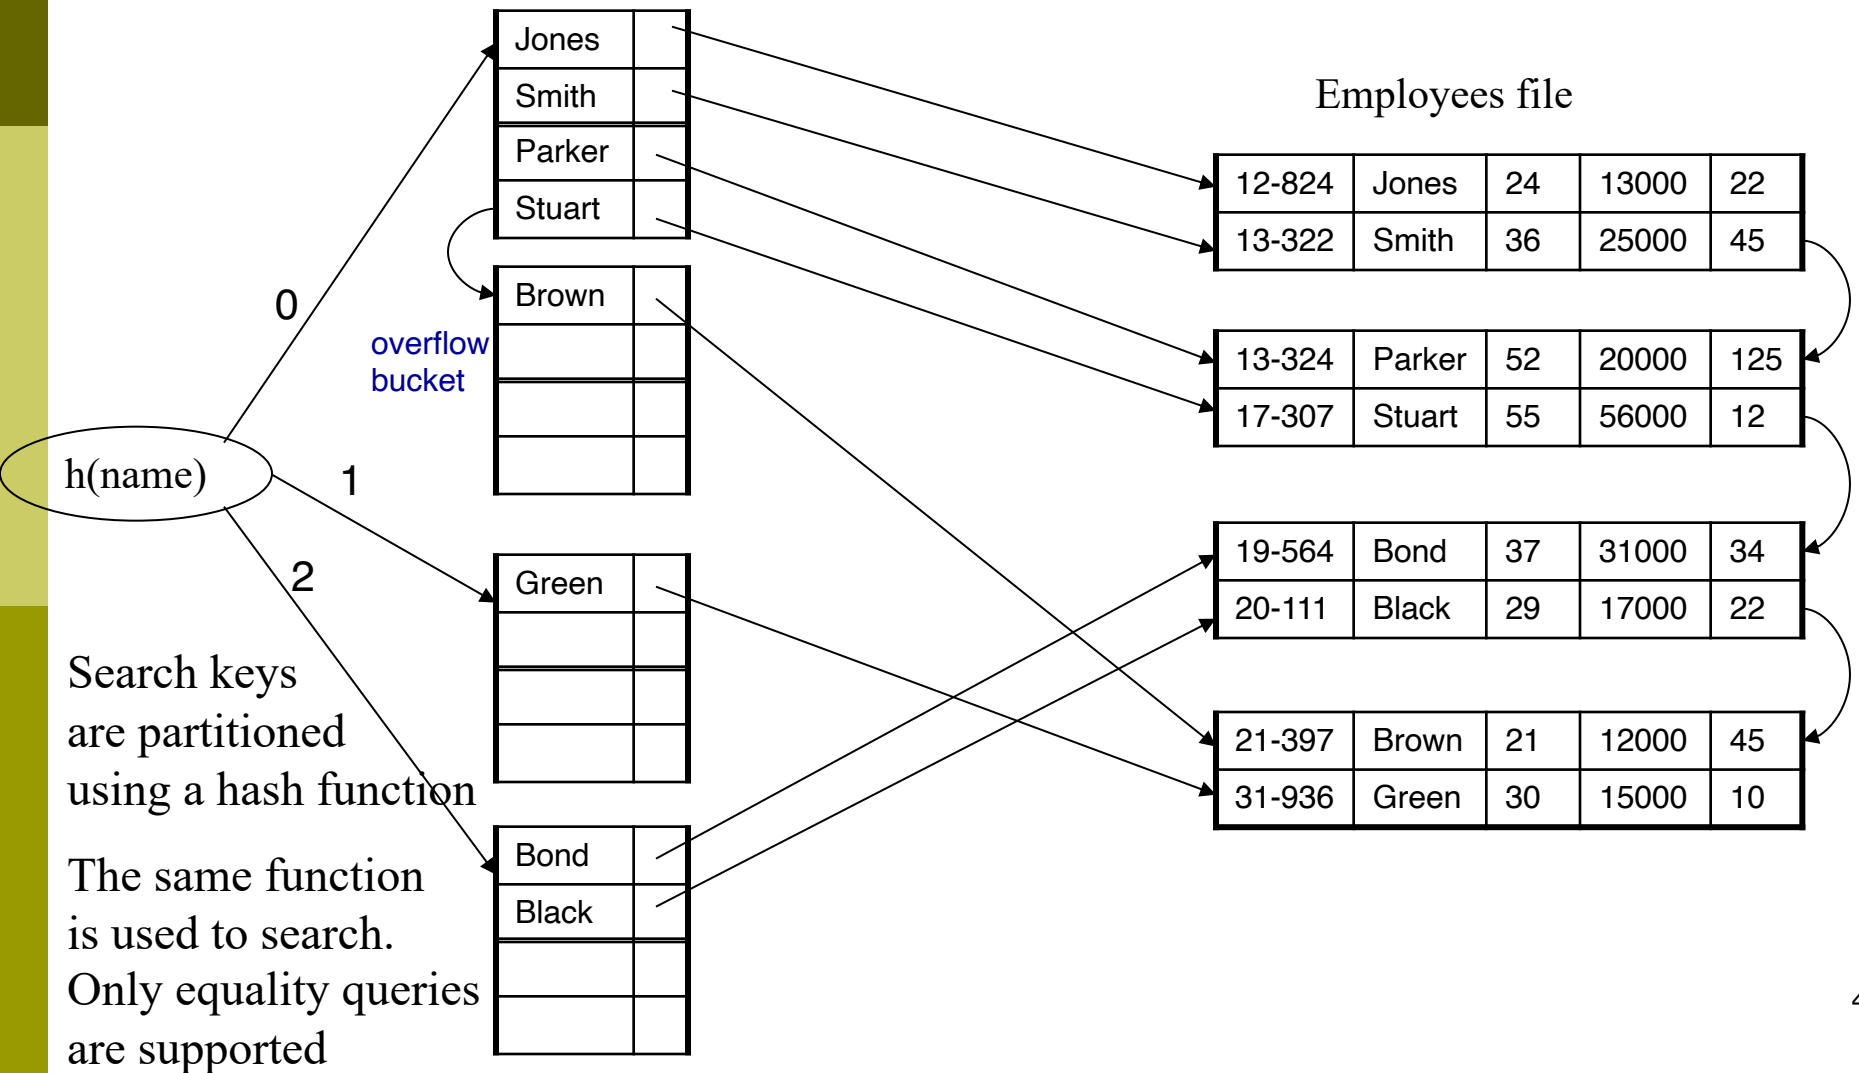
The leaf level of the B⁺-tree contains
<rid, search-key> values sorted by search-key

The non-leaf nodes of the tree are used to guide
search

B⁺-trees are appropriate for

- equality queries
- range queries

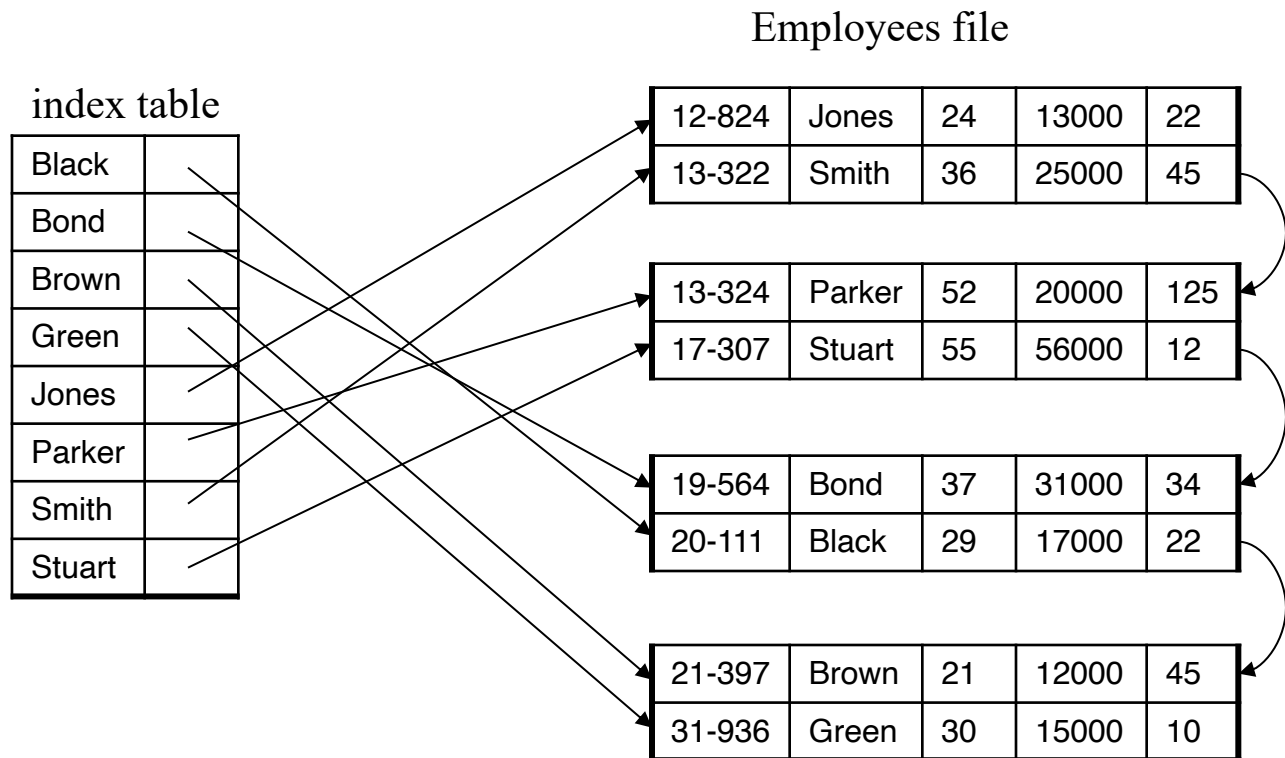
Example of Hash Index



Indexing in main memory

- ❑ B⁺-trees and hash tables can be constructed and used in main memory
- ❑ Motivation: In most applications nowadays data can fit in memory
- ❑ **Sorting** the index table by search-key and using binary search can be an effective index
 - Simple solution that can be used both for equality and range queries
- ❑ **Binary search trees** (AVL, red-black) can also be used if we want to support updates

Sorted index table



When do we create a database index?

- ❑ Indexes are expensive to build and maintain, so they should be used only when necessary
- ❑ Offline analysis
 - Most commercial DBMS include auto-tuning tools
 - Monitor a running workload for a while and then decide what indexes to create or drop based on query patterns
 - Not appropriate if there is no query history or query workload is dynamic
- ❑ Online analysis
 - Monitor workload and performance while processing queries
 - Trigger (dynamically), the creation of new indexes or dropping unnecessary indexes
 - Monitoring and creating/dropping full indexes is quite expensive

Database Cracking



Reading: S. Idreos, M. L. Kersten, and S. Manegold,
“[Database Cracking](#),” in Proceedings of the 3rd International
Conference on Innovative Data Systems Research (CIDR),
Asilomar, California, 2007, pp. 68-78

Adaptive indexing: motivation

- ❑ Constructing a complete index is **expensive** compared to a simple data scan
- ❑ If the index is **not used frequently** its construction does not pay off
- ❑ Some parts of the data may be searched more frequently than others
 - It may be worth to index the data **only partially**
- ❑ **Solution**: database cracking
 - Evaluate queries using linear scan
 - **Opportunistically update index** based on queries, then **progressively guide search**

Database cracking: background

- ❑ DB cracking is based on sorting algorithms such as **quicksort**
- ❑ Quicksort places a given value (the pivot) in the **correct position** in the array
 - Performs **in-place partitioning** of the array into two **pieces**, based on the pivot (smaller than, greater than)
- ❑ Then, it **recursively sorts** the two partitions
- ❑ Example:
 - array: {52, 37, 63, 14, 17, 8, 6, 25}, **pivot**: 25
 - array becomes {6, 8, 17, 14, **25**, 63, 37, 52}
 smaller than pivot greater than pivot

Quicksort algorithm

algorithm quicksort(A, lo, hi)

if $0 \leq \text{lo} \ \&\& \ \text{lo} < \text{hi} \ \&\& \ \text{hi} \leq |A|-1$ **then** *// indices in correct order*
 $p := \text{partition}(A, \text{lo}, \text{hi})$ *// partition A from lo to hi and get pivot position*
 quicksort(A, lo, p) *// recursively sort first partition*
 quicksort(A, p + 1, hi) *// recursively sort second partition*

algorithm partition(A, lo, hi)

$\text{pivot} := A[\text{floor}((\text{hi} + \text{lo}) / 2)]$ *// choose as pivot value the middle element of array*
 $i := \text{lo} - 1$ *// left index*
 $j := \text{hi} + 1$ *// right index*
 loop forever
 do $i := i + 1$ **while** $A[i] < \text{pivot}$
 do $j := j - 1$ **while** $A[j] > \text{pivot}$
 if $i \geq j$ **then return** j *// if indexes cross then pivot is in position j and partitions are correct*
 swap $A[i]$ **with** $A[j]$ *// because $i < j$, $A[i]$ and $A[j]$ are in wrong partitions, so swap them*

Pivot-based partitioning example

$\begin{matrix} i & & & & j \\ \{52, 37, 63, 14, 17, 8, 6, 25\} \end{matrix}$ pivot := 14
lo = 0
hi = 7

$\begin{matrix} i & & & & j \\ \{52, 37, 63, 14, 17, 8, 6, 25\} \end{matrix}$ swap!


$\begin{matrix} i & & & & j \\ \{6, 37, 63, 14, 17, 8, 52, 25\} \end{matrix}$

$\begin{matrix} i & & & & j \\ \{6, 37, 63, 14, 17, 8, 52, 25\} \end{matrix}$ swap!

$\begin{matrix} i & & & & j \\ \{6, 8, 63, 14, 17, 37, 52, 25\} \end{matrix}$

$\begin{matrix} i & & j \\ \{6, 8, 63, 14, 17, 37, 52, 25\} \end{matrix}$ swap!

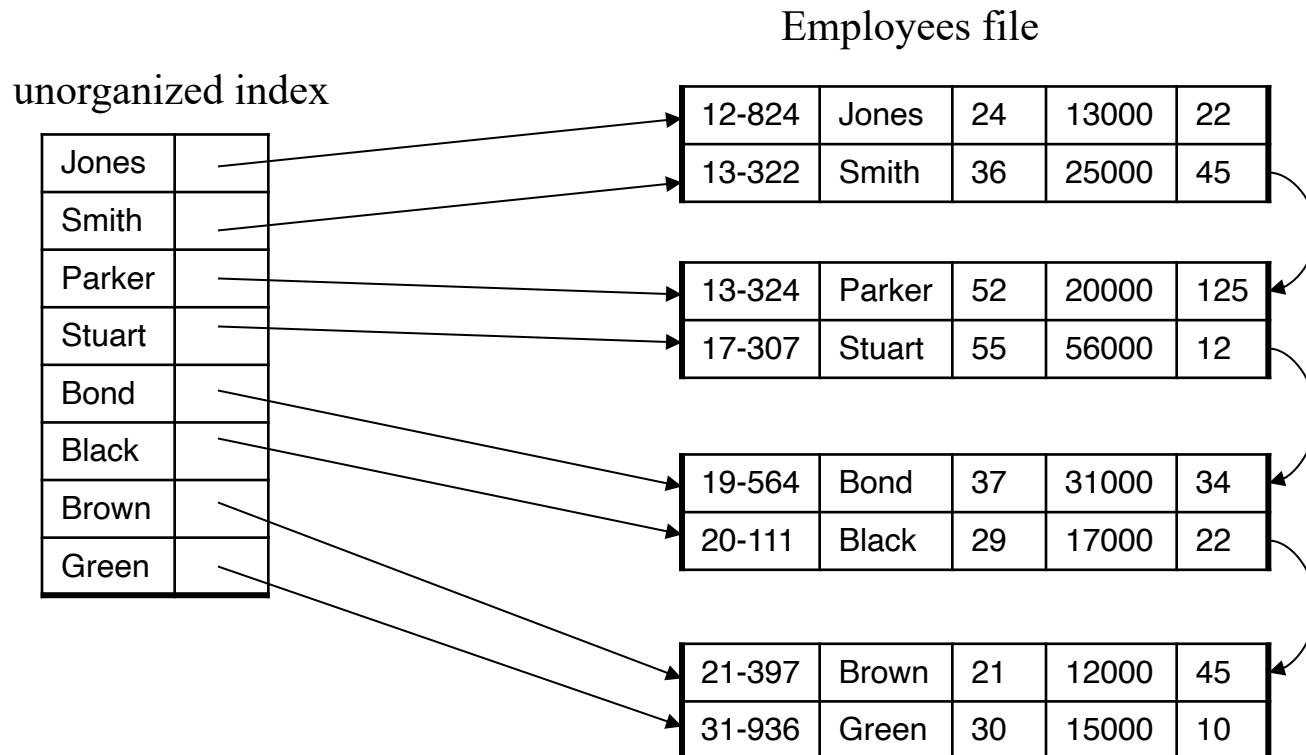
Pivot-based partitioning example



- With one pass over the data, we **crack** the array into two pieces based on pivot value

Standard database cracking

- Initially, we have an unorganized index



- For simplicity, assume that the index is an **unordered array**

Standard database cracking (cont'd)

- Example initial index (keys only):
 $\{52, 37, 63, 14, 17, 8, 6, 25\}$
- Example first query: $10 < \text{key} \leq 20$
- First, crack using $10 < \text{key}$

i \rightarrow

← j

- {52, 37, 63, 14, 17, 8, 6, 25}
- {6, 37, 63, 14, 17, 8, 52, 25}
- {6, 37, 63, 14, 17, 8, 52, 25}
- {6, 8, 63, 14, 17, 37, 52, 25}

- {6, ^j**8**, ⁱ**63**, 14, 17, 37, 52, 25}


Standard database cracking (cont'd)

- Example initial index (keys only):
 $\{52, 37, 63, 14, 17, 8, 6, 25\}$
- Example first query: $10 < \text{key} \leq 20$
- Second, crack using $\text{key} \leq 20$


	$i \rightarrow$		$\leftarrow j$
■	{6, 8, 63, 14, 17, 37, 52, 25}		
■	{6, 8, 63, 14, 17, 37, 52, 25}		
■	{6, 8, 17, 14, 63, 37, 52, 25}		
		j	i
■	{6, 8, 17, 14, 63, 37, 52, 25}		
	$10 < \text{key}$	$\text{key} \leq 20$	

Standard database cracking (cont'd)

- After each query, the partition algorithm of quicksort is used to progressively crack the array
 - Query boundaries define the pivots
- A **binary search tree** organizes the cracks




each node of the BST corresponds to a subarray
e.g. node 10 covers C1C2C3



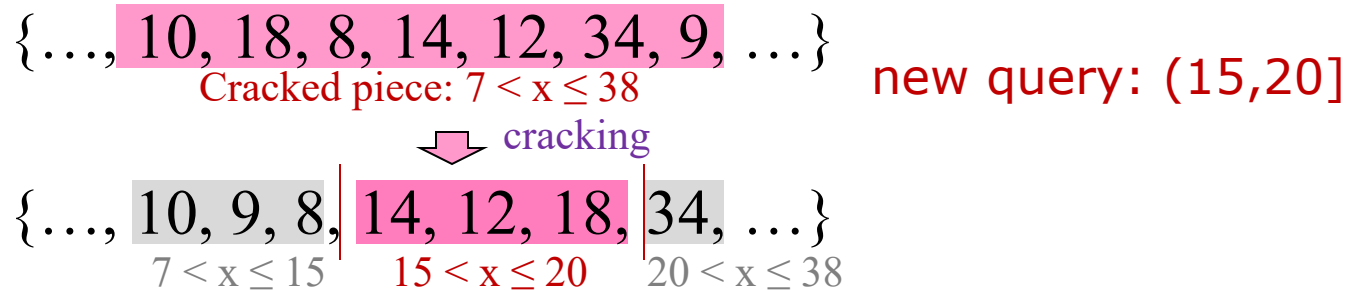
Standard database cracking (cont'd)

- Example: new query (10,12]
 - Search tree (depth-first search)
 - Perform cracks at **leaves** whose range overlaps with query range; collect **query results** from new pieces

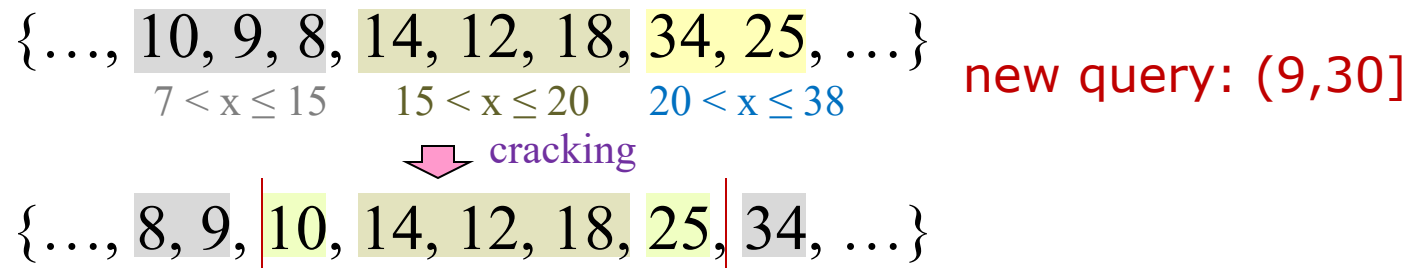


Crack-in-two vs crack-in-three

- Range queries select values in an interval range
- If the interval range falls entirely in an unordered piece the piece should be **cracked in three**




- Otherwise, at most two **cracks-in-two** are applied



Cumulative query cost

Data: unordered array with 10 million distinct integers

Queries: $v_1 < x < v_2$, where v_1 and v_2 are random



Simple: answer each query by linear scan

Sort: sort data before 1st query, then binary search


Cracking: perform cracking during queries

Notes on standard cracking

- ❑ After large number of random queries, the array becomes completely sorted
 - In practice, a piece is not cracked if it is very small
 - i.e. when sequentially scanning the piece is faster than BST-based searching of the corresponding subarray
- ❑ If the query workload is skewed and most queries fall within a limited range, we avoid accessing and/or sorting large parts of the array
- ❑ Cracking can be used to adaptively index columns in a column-store
- ❑ Cracking is not appropriate for disk-based data

Best case for cracking


- Each query is a sub-range of the previous one



- Cracked pieces are always small

Worst case for cracking

- Each query is a super-range of the previous one



- Big pieces are cracked multiple times!

Worst case for cracking

- First query asks for smallest element, second query for second smallest etc.



- Big pieces are cracked multiple times!

Adaptive Merging



Reading: Goetz Graefe, Harumi A. Kuno: Self-selecting, self-tuning, incrementally optimized indexes. EDBT 2010: 371-381

Adaptive Merging

- ❑ Motivated by merge-sort, instead of quick-sort
- ❑ Initially, divides the data into small partitions and sorts each of them
 - Sorting small partitions is much faster than sorting the entire array
- ❑ Each query selects segments from sorted partitions and merges them into a new partition
 - Binary search is used to search in partitions
- ❑ Merging may cause some original partitions to disappear
 - Eventually everything is merged to a single sorted array

Adaptive Merging (example)

{52, 37, 63, 14, 17, 8, 6, 25, 7, 10, 38, 20} original array

⇓ first pass: create initial sorted partitions

{14, 37, 52, 63} {6, 8, 17, 25} {7, 10, 20, 38}

⇓ search each partition and merge

new query: (15,20]

{17, 20} {14, 37, 52, 63} {6, 8, -, 25} {7, 10, -, 38}

⇓ search each partition and merge

new query: (5,10]

{6, 8, 7, 10} {17, 20} {14, 37, 52, 63} {-, -, -, 25} {-, -, -, 38}

↖ ↗ not touched (out of query bounds)

Notes on Adaptive Merging

- ❑ B-tree-like indexes can be used to accelerate search in each partition
- ❑ Multi-way merge can be used to merge multiple segments
- ❑ Results of overlapping queries are stored in single partitions
 - query results are also partitions, which are **merged** with next queries
 - partitions that are query results are never split/cracked
 - eventually, **all data are merged to a single sorted array** with the results of all queries

Adaptive Merging vs. DB cracking

- ❑ **Slow startup**, as it requires sorting each initial partition
- ❑ Creating new partitions allocates new memory space and moves data in memory
 - “gaps” are created in original partitions
 - **not as efficient as swaps performed by cracking**
- ❑ **Harder to implement compared to DB cracking**
- ❑ Partitions that are query results are not split
 - excessive fragmentation is avoided
- ❑ **Avoids moving a key multiple times**, which may happen often in DB cracking
- ❑ **Converges to a single sorted array much faster compared to DB cracking**

Hybrid Adaptive Indexing

- ❑ **Main drawback of DB cracking**: slow convergence
- ❑ **Main drawback of adaptive merging**: slow startup
- ❑ **Objective**: a hybrid approach that combines the best from both strategies
- ❑ **Main idea**: start with some initial partitions as in adaptive merging, but not necessarily sorted
 - Option 1: apply **cracking** to initial partitions, to obtain the segments that should be merged to form a query result
 - Option 2: apply **radix-clustering** to initial partitions
 - ❑ Convert values to integer codes, use k most significant bits of codes to divide partition to 2^k segments
- ❑ **Option 1 performs well in an experimental analysis**

Hybrid Adaptive Indexing(example)

{52, 37, 63, 14, 17, 8, 6, 25, 7, 10, 38, 20} original array

⇓ create initial partitions (not sorted)

{52, 37, 63, 14} {17, 8, 6, 25} {7, 10, 38, 20}

⇓ crack each partition and merge

new query: (15,20]

{17, 20} {14, 37, 63, 52} {8, 6, -, 25} {7, 10, -, 38}

⇓ crack each partition and merge

new query: (5,10]

{6, 8, 7, 10} {17, 20} {14, 37, 52, 63} {-, -, -, 25} {-, -, -, 38}

↖ ↗ not touched (out of query bounds)

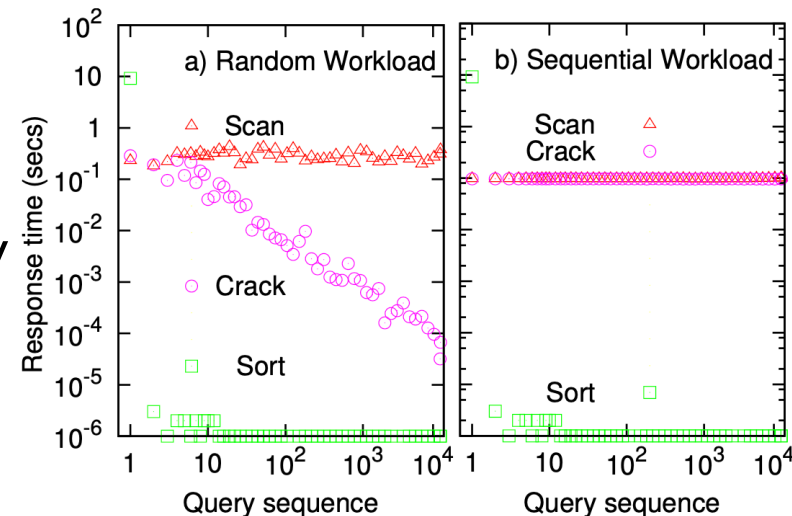
Stochastic Cracking



Reading: F. Halim, S. Idreos, P. Karras, and R. H. C. Yap, "[Stochastic Database Cracking: Towards Robust Adaptive Indexing in Main-Memory Column-Stores](#)," Proceedings of the Very Large Databases Endowment (PVLDB), vol. 5, no. 6, pp. 502-513, 2012

Stochastic Cracking

- **Motivation:** standard cracking assumes a random query workload
 - No prediction or care for future queries
 - Large pieces are left completely unindexed
 - Performs badly for specific workloads (**sequential**)
- However, access pattern is typically unknown in exploratory analysis
- **Main idea:** proactively perform additional cracks that can serve future queries
 - Extra crack points are determined to some degree by the queries



Stochastic Cracking: DDC

- Data Driven Center strategy
- Given a query $Q=[a,b]$
 - Recursively **halve** the piece(s) where Q falls, until it (they) become(s) sufficiently small
 - Then, crack the piece(s) based on $[a,b]$

{52, 37, 63, 14, 17, 8, 6, 25, 7, 10, 38, 20} initial array

↓ first pass: crack based on median

new query: (6,7]

{14, 17, 8, 6, 7, 10} {52, 37, 63, 25, 38, 20}

↓ second pass on 1st piece: crack based on median

{8, 6, 7} {14, 17, 10} {52, 37, 63, 25, 38, 20}

↓ third pass on 1st/1st piece: crack based on query

{6} {7} {8} {14, 17, 10} {52, 37, 63, 25, 38, 20}

Finding median is combined with cracking (in linear time)


Stochastic Cracking: DDR

- Data Driven Random strategy
- Given a query $Q=[a,b]$
 - Recursively **randomly crack** piece(s) where Q falls, until it (they) become(s) sufficiently small
 - Then, crack the piece(s) based on $[a,b]$
- Uses **randomly picked pivots** for cracking recursively
- Avoids the overhead of finding medians
- Remains query-driven, since it recursively cracks only the piece that overlaps with query range

Stochastic Cracking: Other heuristics

- ❑ **DD1C** same as DDC, but performs at most one center crack before final query crack
- ❑ **DD1R** same as DDR, but performs at most one random crack before final query crack
- ❑ **MDD1R** same as DD1R, but does not perform final query crack; instead it copies query results to a separate array
 - Result identification and copying is done at the same time as the random crack
- ❑ **Progressive MDD1R**: perform only a small percentage of swaps at each crack
 - Lowers burden of first few queries
- ❑ **Selective Stochastic Cracking**:
 - selectively eschew stochastic cracking for some queries
 - such queries are answered using original cracking

Stochastic Cracking: all approaches



MDD1R is experimentally shown to be the best option in terms of overall cost for different query workloads

Coarse Granular Index



Reading: Felix Martin Schuhknecht, Alekh Jindal, Jens Dittrich:
[An experimental evaluation and analysis of database cracking.](#)
VLDB J. 25(1): 27-52 (2016)

Coarse Granular Index

- **Motivation:** First queries in DB cracking are slow because large parts of the data are read and possibly swapped
- **Idea:** build a **coarse granular index**
 - Create **balanced partitions** using range partitioning upfront
 - All data in one partition are smaller than all data in the next partition
 - Apply standard cracking later on
 - The cost of the first query is higher than standard cracking but significantly lower than full indexing

Coarse Granular Index

{52, 37, 63, 14, 17, 8, 6, 25, 7, 10, 38, 20} initial array

↓ first query: partition data + crack

new query: (6,7]

{14, 17, 8, 6, 7, 10, 20} {37, 25, 38} {52, 63}

{6} {7} {14, 17, 8, 10, 20} {37, 25, 38} {52, 63}

↓ next queries: identify partition(s) + crack



new query: (15,20]

{6} {7} {8, 10, 14} {17, 20} {37, 25, 38} {52, 63}

Coarse Granular Index

Experiment:

- 10^8 uniformly distributed values with a key range of $[0; 100,000]$
- 1000 random queries of the form $\text{low} \leq x < \text{high}$, each with selectivity 1%



Multidimensional Adaptive Indexing




Reading: Matheus Agio Nerone, Pedro Holanda, Eduardo Cunha de Almeida, Stefan Manegold: Multidimensional Adaptive & Progressive Indexes. ICDE 2021: 624-635

Adaptive KD-tree


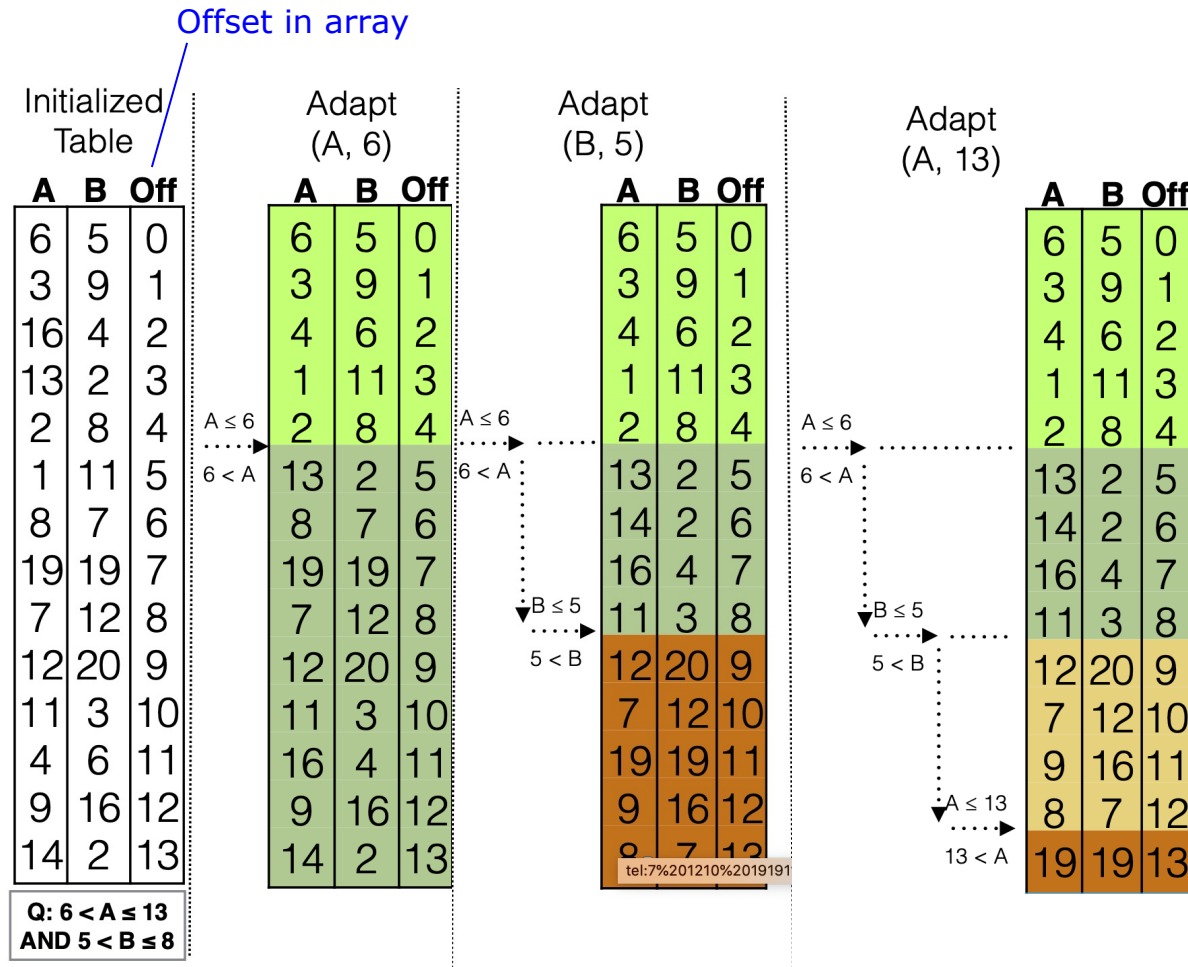
- ❑ **Target:** Exploratory range query workloads in multidimensional spaces
- ❑ **Idea:** Apply cracking in a multidimensional space using a KD-tree to index the cracked pieces
 - Cracking is performed at one dimension at a time
 - Each node of the KD-tree holds one attribute (dimension) and a value (split point)
 - Leaf nodes point to cracked pieces

Background: the KD-tree

- ❑ Reading: https://en.wikipedia.org/wiki/K-d_tree
- ❑ A **binary tree** for points in a k-dimensional space
- ❑ Root uses one point (median) to divide the space into two subspaces using one dimension
 - ❑ sort a sample to find median
- ❑ Each node uses a point to divide its subspace in two parts
- ❑ Dimensions alternate per level



Adaptive KD-tree example



Cracking not
done because
size is 4 or less



first crack using lower Q bounds in each dim, then upper Q bounds
size threshold for cracking = 4

Notes on Adaptive KD-tree

- Progressively cracks each piece that overlap with query range
 - Crack based on lower bound in dimension 1, then based on l.b. in dim. 2, and so on
 - Repeat the above using upper query bounds in each dimension
 - Stop cracking when the piece becomes smaller than or equal to size threshold
 - At each crack a new KD-tree node is defined

Index Lookup

- When a new query arrives, the current KD-tree is used to guide search
 - When a leaf node is reached, it is cracked if necessary to introduce new nodes



Leaf node reached; since size is 4 or less, it is not cracked

new query:

Q: $6 < A \leq 15$
AND $0 < B \leq 5$

Summary

- ❑ Adaptive, self-organizing indexes are dynamically constructed as a result of query evaluation
- ❑ Database cracking was the first idea toward this direction
- ❑ Another approach is adaptive merging
- ❑ Stochastic cracking improves standard cracking to handle non-random query workloads
- ❑ A cheap and coarse data partitioning before cracking is beneficial
- ❑ Cracking also applied to multidimensional data
 - KD-tree can be used to index cracked pieces