

# Learned Indexes

---

- ❑ Motivation: use ML models in place of data structures and algorithmic tasks
- ❑ A general framework for learned indexes
- ❑ A practical and updateable learned index
- ❑ Multidimensional learned indexes

# Learned Indexing Framework



**Reading:** Tim Kraska, Alex Beutel, Ed H. Chi, Jeffrey Dean, Neoklis Polyzotis: [The Case for Learned Index Structures](#). SIGMOD Conference 2018: 489-504

# Learned Indexing: Motivation

---

- **General trend**: replace and augment traditional algorithms with machine learning (ML) models
  - e.g. recommender systems
- **Key idea**: Indexes are “models” that predict the position of a key in a dataset


# An Extreme Example

---

- Suppose that we want to index an array of 100M consecutive integers

0 1 2 3 4 5 6 7 ... 999999999

- We would not use a B+-tree, because the position of a number implies its value!
- This is an extreme example of using a model to index, i.e., a **learned index**



- Search cost:  $O(1)$ , space complexity:  $O(1)$


# Practical Issues

---


- ❑ The data do not always have a distribution that can be captured by a single model
- ❑ The effort to learn the data distribution can be high
- ❑ The data distribution may change
  - dynamic data
- ❑ **Goal of Learned Indexing**: Use ML to define models (e.g. neural networks) that can enhance or replace, traditional index structures (e.g. B<sup>+</sup>-Trees, hash indexes, Bloom filters)

# Black-box vs. Glass-box Models

---



Deep Neural Networks  
(black-box)



Linear regression model  
(glass-box)

# Existing Indexes as Models

---

## □ B<sup>+</sup>-tree:

- The sequence of leaf nodes can be thought of as a sorted array
- Non-leaf nodes can be thought of as models that predict the position of a key in the array

## □ Bloom filter:

- Think of it as a binary classifier: predicts if a key exists in a set or not


## □ Important differences:

- B<sup>+</sup>-tree does not have errors
- Bloom filter can have false positives but no false negatives


# Learned Index for Range Searching

- Based on the B<sup>+</sup>-tree
- Assumes static data, where  $\langle \text{key}, \text{rid} \rangle$  pairs are sorted in an array
  - Leaf nodes of the tree are pieces in the array
- Replaces inner tree nodes by model(s)

(a) B-Tree Index



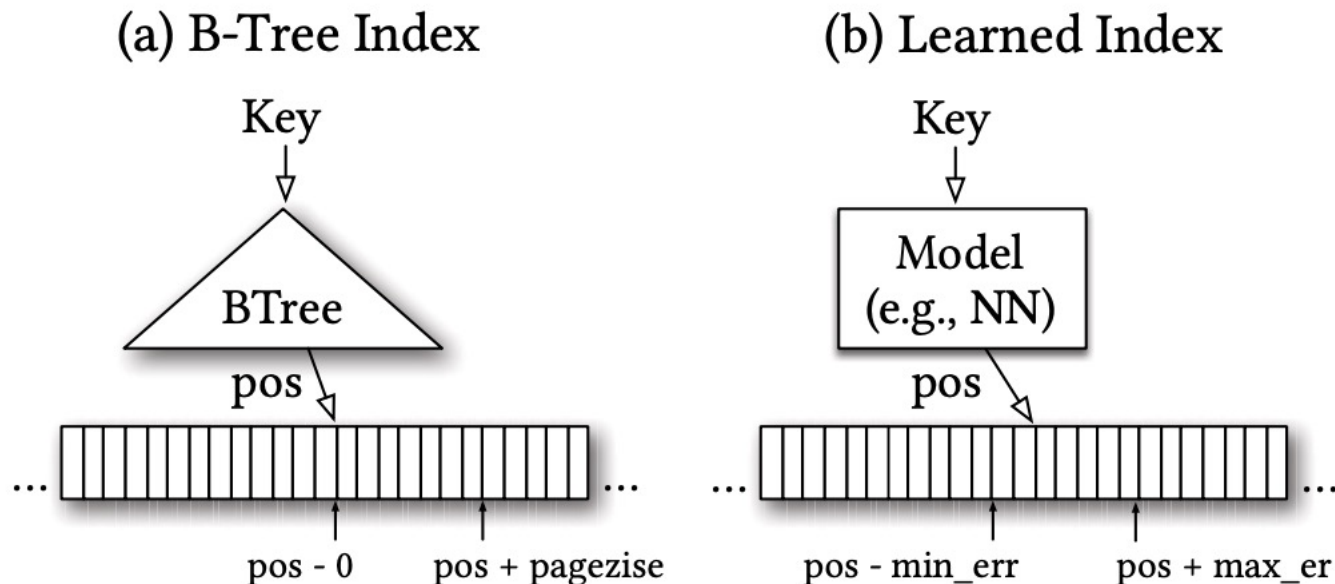
(b) Learned Index





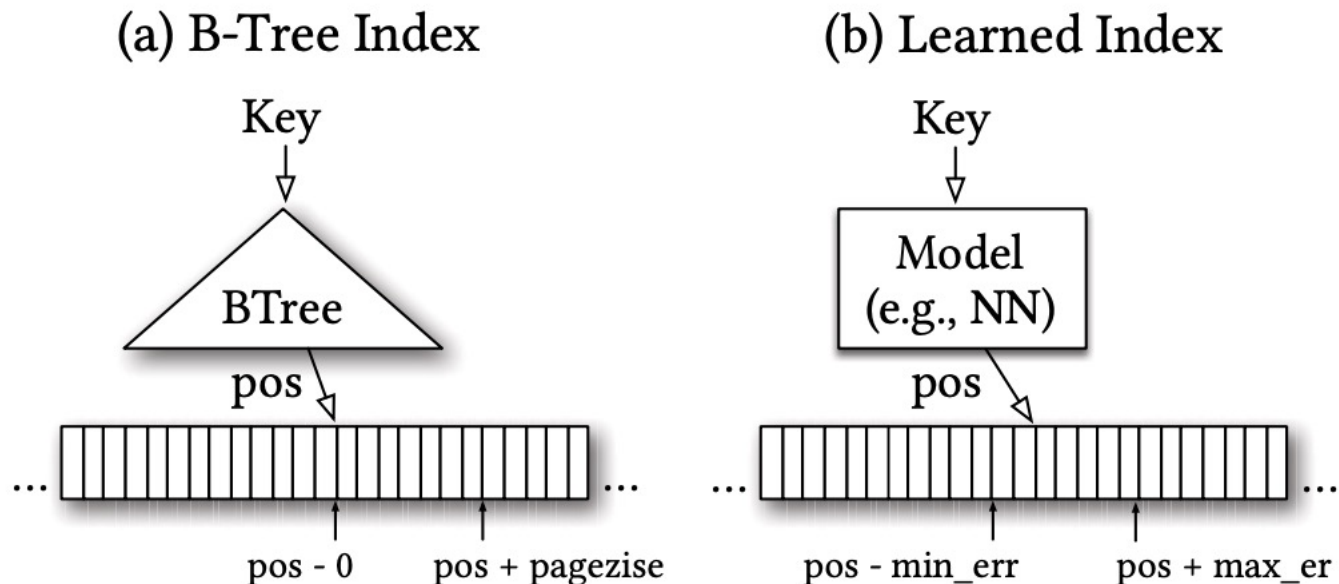
# Learned Index for Range Searching

- The B+-tree is a **regression tree** model
  - maps a key to a position with a  $\text{min\_error}=0$  and  $\text{max\_error}=\text{page\_size}$
  - guarantees that the key can be found in that region if it exists



# Learned Index for Range Searching

- ❑ No need to have a fixed error bound
  - any error easily corrected by local search around the prediction (e.g., [exponential search](#))
- ❑ Exploit modern hardware and ML accelerators




# Range Index Models as CDF Models

- **Model input:** key
- **Model output:** position of key in the array
- **Observation:** a model for the position of a key in a sorted array approximates the cumulative distribution function (CDF)

$$p = F(\text{key}) * N$$

key position estimate      CDF  $P(x \leq \text{key})$       number of records




# Range Index Models as CDF Models

---

- Possible solution: Model the CDF
  - Train a linear regression model to learn the CDF
  - Minimizing the (squared) error of a linear function
  - Alternative: use a neural net
- Using a single model for the whole CDF will not work
  - Slow training, poor “last mile” accuracy

# Recursive Model (RM) Index

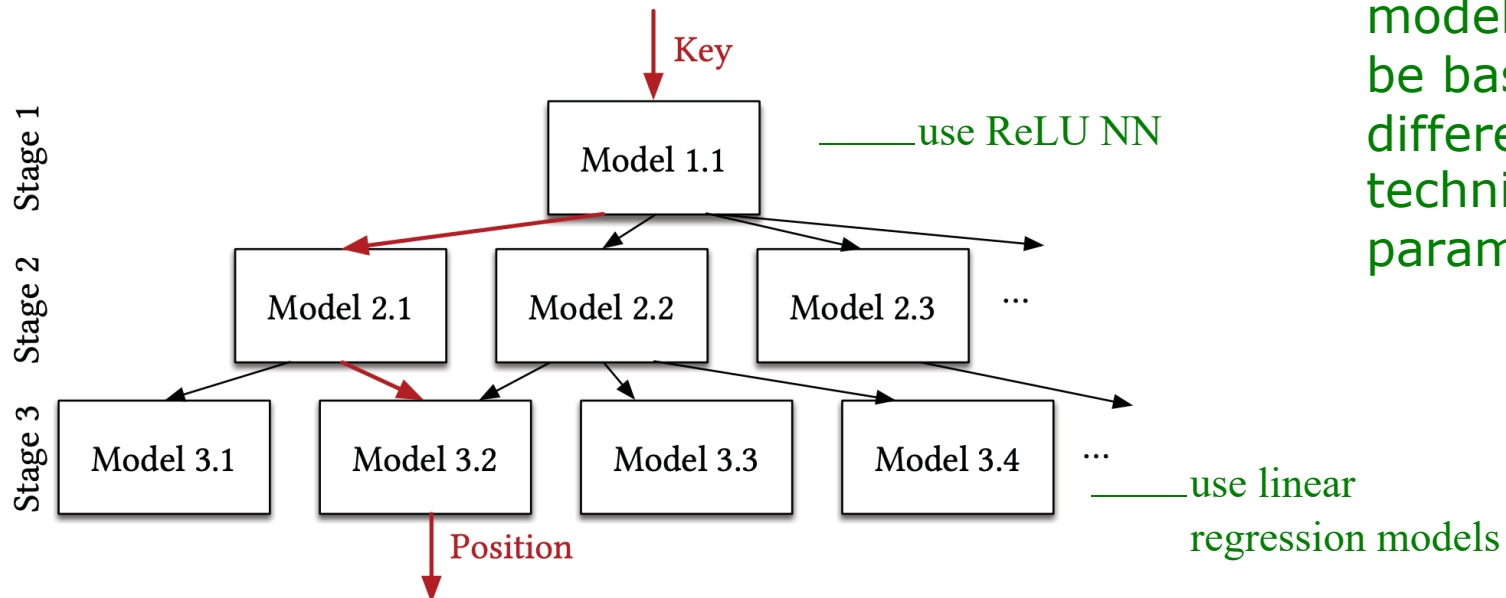
- A hierarchy of models
  - A non-leaf node in the hierarchy is a model that predicts which child model to use,
  - A leaf node is a model to predict the position of the key in an array



# Recursive Model (RM) Index

- At each stage, a model takes the key as input and **selects** the model to use at the next level
  - Models need not form a tree
- Use **local search** at the “last mile”

Different models can be based on different techniques/parameters



keep min- and max-error for every leaf model

# Learned Index for Value Searching

---

- **Hash-maps** are effective for single-value search, but suffer from collisions
  - Hard to avoid bucket overflows
- ML has been used to find hash functions that avoid collisions
- **Idea**: Learning the CDF of keys can help to learn a good hash function
  - Scale the CDF by the targeted size  $M$  of the hash-map
  - Use  $h(\text{key}) = F(\text{key}) * M$  as hash function
  - Use the RM-index for  $F$
  - If CDF learning is accurate, collisions are avoided

# Learned Index for Existence Search


---

- ❑ **Search problem:** test whether an element is a member of a set  $S$
- ❑ **Classic solution:** Bloom filter
  - Use a bitarray of size  $m$
  - Use  $k$  **independent hash functions** that map each key to one of the  $m$  positions
  - Use a single bitmap  $B$  for  $S$ ; for each key in  $S$ , set  $B[h_i(s)] = 1$ , for each  $i = 1, \dots, k$
  - For a search key  $x$ , apply all hash functions and get  $x$ 's bitmap  $B_x$
  - If  $B_x$  has 0 in any position where  $B$  has 1,  $x$  is not a member of  $S$
  - Bloom filter guarantees there are **no false negatives**, but **there may be false positives** ( $x$  mistaken as a member)



# Learned Index for Existence Search

- ❑ **Issue**: a good Bloom filter takes too much space
- ❑ **Solution**: Learned Bloom filter
  - **Goal**: learn a good hash function with **lots of collisions among keys** and **lots of collisions among non-keys**, but **few collisions of keys and non-keys**
  - Training needs not only  $S$ , but also a set  $U$  of **non-keys**
    - ❑  $U$  obtained from past queries or could be random
  - Possible models: binary classifiers (e.g. RNN or CNN)
    - ❑  $D = \{(x_i, y_i = 1) | x_i \in S\} \cup \{(x_i, y_i = 0) | x_i \in U\}$
  - Choose a threshold  $\tau$ ; assume that  $x$  exists if  $f(x) \geq \tau$
  - Create an **"overflow" Bloom filter** for false negatives from  $f$



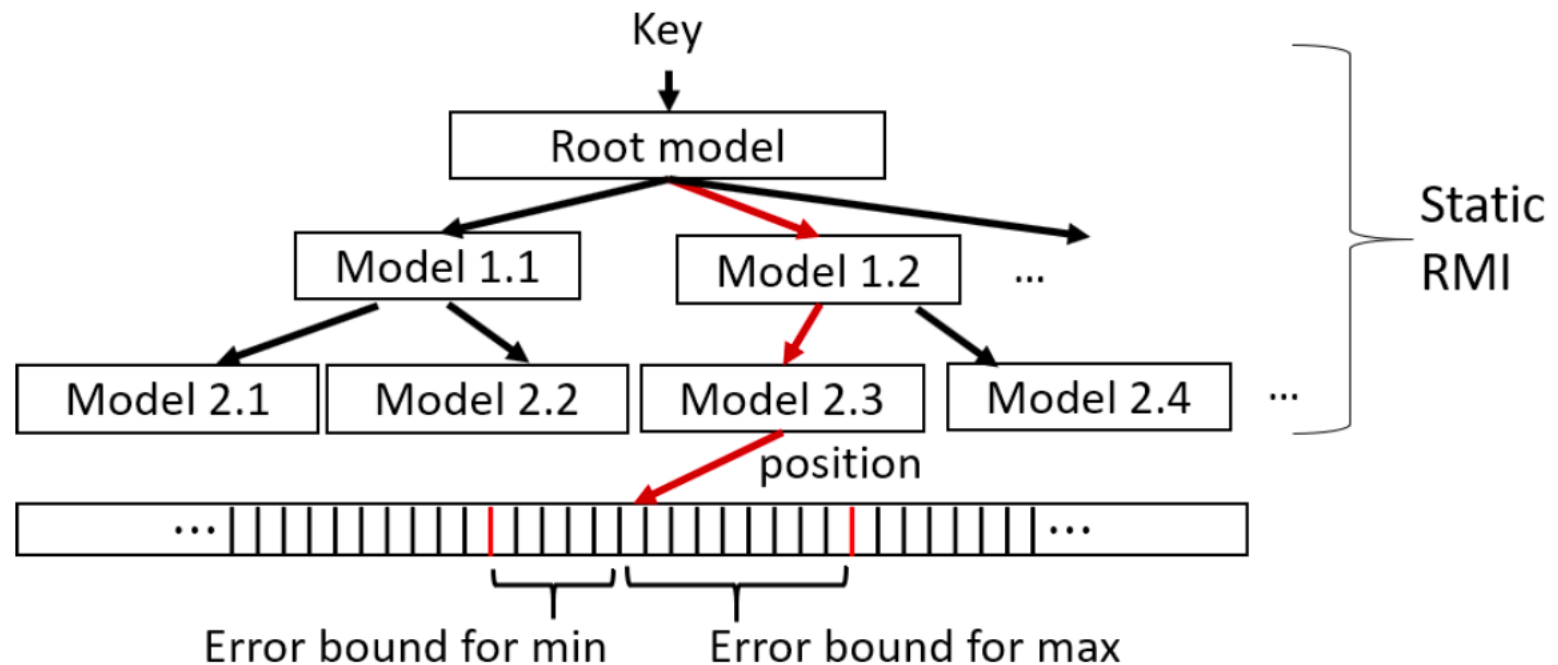
# Adaptive Learned Index



**Reading:** J. Ding, U. F. Minhas, J. Yu, C. Wang, J. Do, Y. Li, H. Zhang, B. Chandramouli, J. Gehrke, D. Kossmann, D. B. Lomet, T. Kraska: [ALEX: An Updatable Adaptive Learned Index](#). SIGMOD Conference 2020: 969-984

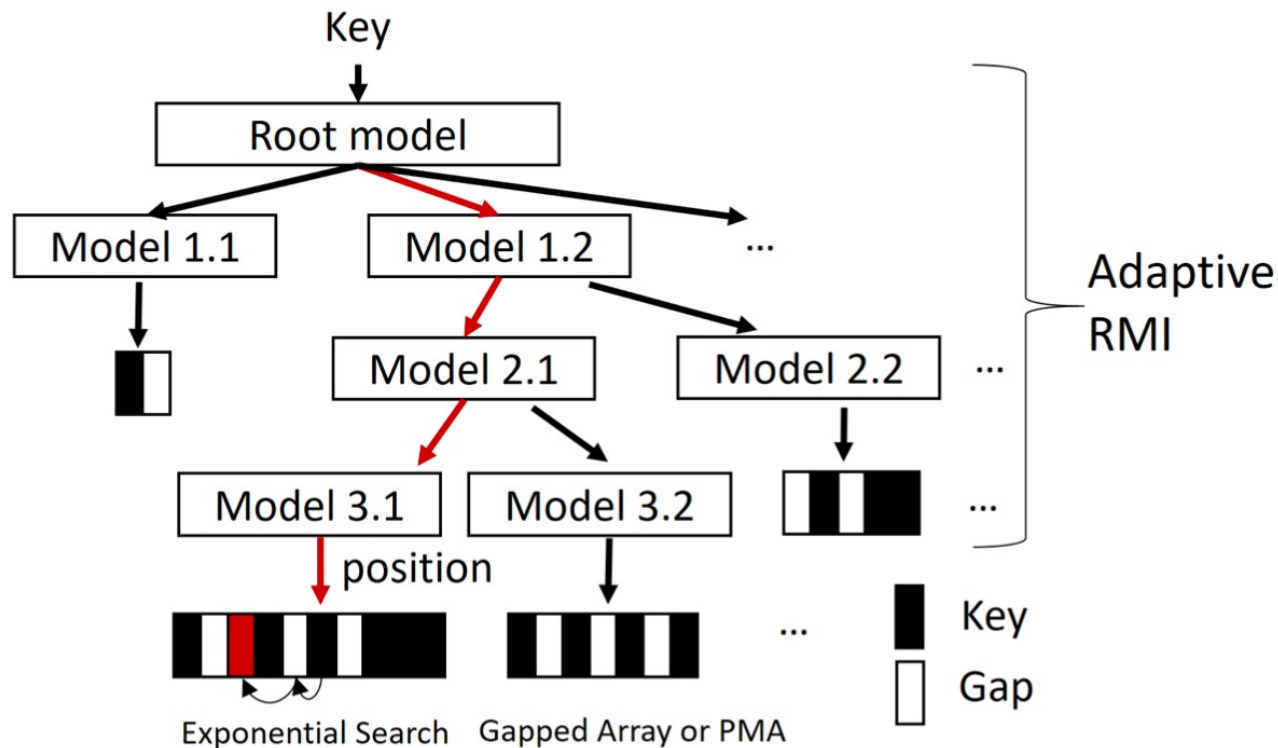
# Drawback of RM-index

- ❑ **Static depth** and **static models** in the index
  - Models become ineffective as data are updated
- ❑ The keys are in a sorted **static array**
  - Updates are expensive



# ALEX: Adaptive Learned Index

- Primary goal: support updates
- Differences to RM-index:
  - Dynamic leaves of different structure
  - Adaptive RM-index based on the workload



# Performance goals of ALEX

---

- ❑ Insert time should be competitive with B<sup>+</sup>-tree
- ❑ Search time should be faster than B<sup>+</sup>-tree and similar to RM-index
- ❑ Index storage space should be similar to RM-index and much smaller than B<sup>+</sup>-Tree
- ❑ Data storage space (leaf level) should be comparable to B<sup>+</sup>-Tree
  - Low index storage space is important in cases where index storage is in (limited) memory and data storage on disk

# Features of ALEX

---

- ❑ Use **exponential search** instead of binary search
- ❑ Use **dedicated arrays for leaf nodes** instead of a single sorted data array
- ❑ Use **gaps in dedicated arrays** to facilitate fast insertion
- ❑ **Avoids** partitions with very few keys and partitions fully-packed with keys
- ❑ Models of non-leaf nodes **adapt** as data are updated


# Exponential vs. Binary Search

---

- ❑ ALEX uses **exponential search without error bounds** to find keys at the leaf level
  - RM-index uses **binary search with error bounds**
- ❑ **Rationale**: if the models are good, their prediction is close enough to the final position, so exponential search will be efficient
- ❑ ALEX does not need to store error bounds for RMI models, which saves space

## Exponential search example:

search for 37, prediction at position of 14



# Leaf Nodes vs. Single Array

---

- ❑ ALEX uses a dedicated array for each leaf
  - RM-index keeps **all keys in a single sorted array**
  - **Insertions are very expensive**
- ❑ Each leaf node is an **array with gaps** between elements to facilitate cheap insertions
- ❑ Two alternative array layouts for leaves:  
**Gapped Array vs. Packed Memory Array**



# Gapped Array Leaf Node Layout

- Allow “naturally” distributed gaps between array elements

$$\text{predicted\_position} = \text{key} / 2$$



- To insert a new key



- use the RMI to predict the insertion position
- if a gap and inserting there keeps the sorted order, all done
- else use exponential search to find correct position
  - If not a gap, make a gap by shifting keys and insert there
- $O(\log n)$  insertion time with high probability
- If  $\text{numkeys}/\text{leafsize} \geq \text{density bound } d$ 
  - expand leafsize by  $1/d$ , re-train model for leaf
  - (must keep enough gaps to prevent shifting)
- **Drawback of gapped array:** a possible long continuous region with no gaps (expensive shifting)

# Packed Memory Array (PMA) Layout

- Node size is a power of 2 and is divided to segments whose size is also a power of 2



- An **implicit binary tree** on top of the array
  - Each node has a **density bound** for non-gaps
  - Nodes near leaves have higher density bounds
  - If insertion to a segment violates bound, redistribute keys to ensure no bound is violated
  - If no redistribution solves problem, **double size of array**
- Redistribution may affect model's accuracy**
  - Perform model-based redistribution
  - Retrain model after **node expansion**

# Adaptive RM-Index


---

- ❑ ALEX tries to avoid two problems
  - **Wasted models**: leaf nodes with very few keys
  - **Fully-packed leaves**: expensive insertions
- ❑ RM-index initialization in ALEX
  - **Goal**: adaptively determine RMI depth and leaves
  - **Recursive building**:
    - ❑ Train a model for root, evenly partition keys to root's children
    - ❑ Oversized partitions (in number of keys) are recursively split
    - ❑ Each non-root node has a fixed number of partitions
    - ❑ Leaves with very few keys are progressively merged with siblings

# Adaptive RM-Index (cont'd)

## □ RM-index updates in ALEX

- **Goal:** adapt RMI to changes
- **Node splits:** if a leaf becomes oversized, it is split to a number of children (new leaves)
  - The model for the split leaf becomes the model for the new parent
  - Data from split leaf node distributed to the new (children) leaves
  - New models are trained for the created children



# Multidimensional Learned Indexes



**Reading:** Vikram Nathan, Jialin Ding, Mohammad Alizadeh, Tim Kraska: [Learning Multi-Dimensional Indexes](#).  
SIGMOD Conference 2020: 985-1000

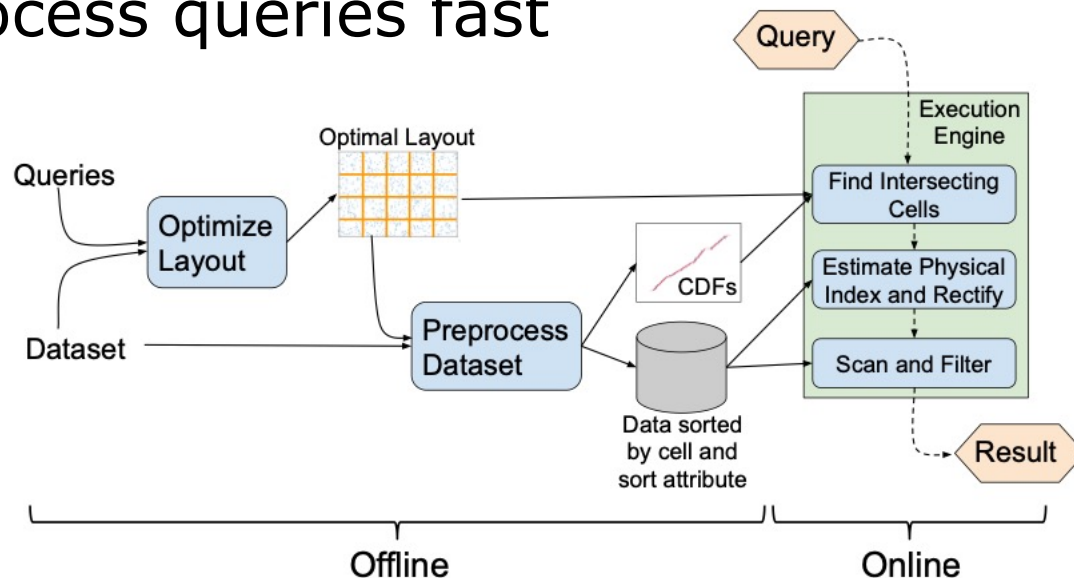
# Multidimensional Indexes

---

- ❑ **Objective**: index multiple attributes simultaneously
- ❑ **Challenge**: data cannot be totally ordered
- ❑ There is room for improving existing multidimensional indexes using ML models, such as in the 1D case
- ❑ **Possible solution**: define a 1D order of the data (e.g. z-order) and then use a 1D learned index
  - **Issue**: z-order CDFs are hard to learn

# A Learned Multidimensional Index

- **Flood** is a multi-dimensional learned index that adapts to the query workload
  - Designed mostly for static data
  - Suitable for multidimensional range queries
  - Created in a preprocessing phase, then used to process queries fast




# Flood: preprocessing

---

- Define an order on the  $d$  dimensions
- Use the first  $d-1$  dimensions to define a  $(d-1)$ -dimensional grid on the data space
  - Dimension  $i$  is divided to  $c_i$  equi-sized partitions
  - ML model is used to determine partitioning
- Each data point is mapped to a grid cell
  - Point coordinates and ML models are used to locate cell
- Cells are sorted in based on a dimension ordering
- In each cell, points are sorted based on  $d$ -th dimension
- Flood uses ML and query workload to find the best dimension ordering and grid granularity




# Flood: 2D data example



- Dimension order ( $x, y$ ) and  $c_0 = 5$
- Points are partitioned to columns along  $x$  and then sorted by their  $y$ -values
- The serialization order is indicated by the arrows

# Flood: query processing

- Find cells that overlap with multidimensional query range, using the ML models
- Use point ordering in each cell to confine the range of points to be scanned
  - Done using a **piecewise linear model**
- Scan range and filter



(1a) Projection finds 4 intersecting cells (cells 1-4) → (1b) Identify physical index range of third cell. (Repeat for other cells.) → (2) Refine the physical index range → (3) Scan and Filter

# Summary

---

- ❑ Learned indexes aim at reducing the memory footprint of indexing while also enabling fast data accesses
- ❑ Performance depends on how well the ML models fit the data distribution
- ❑ Dynamic modeling and gapped array storage is proposed for adaptive indexing
- ❑ Multidimensional learned indexes partition data, then create a 1D sort order at each partition